

hands-on perf workshop

let us into your performance #BCNJUG

May 25 2017, Barcelona JUG

@srvaroa
@duarte_nunes
@gontanon

hello

agenda

Intro & mechanics

The workshop
Exercises

Wrap up

first of all

Check out our repository at:

<https://github.com/srvaroa/jug-perf-workshop>

WIFI details

NR-GUEST - RubyOnRails!

intro and mechanics

Some exercises to let you discover & play with tools

One presenter per exercise.

Presenters will wander around the room supporting

exercise mechanics

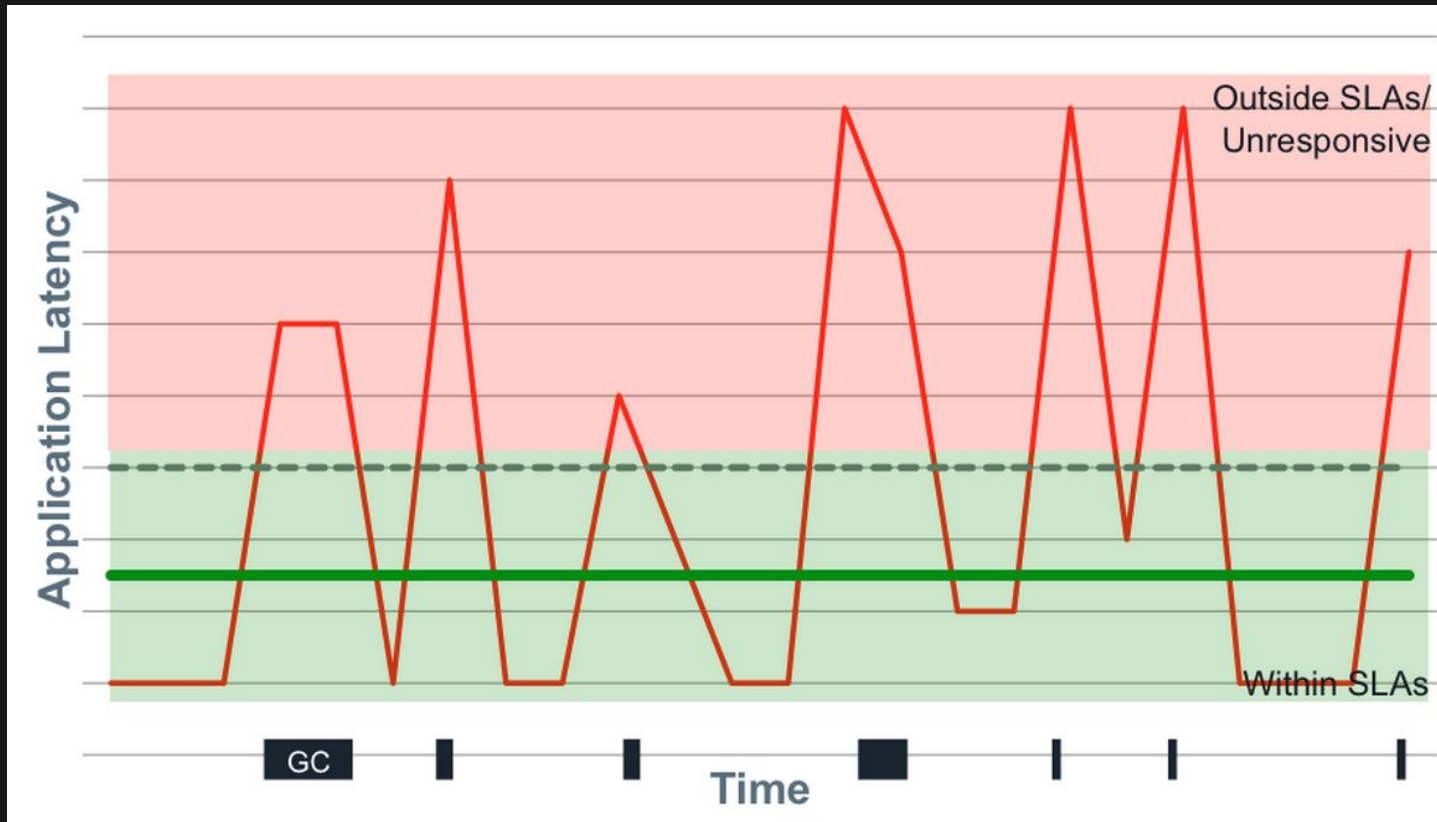
1. Each exercise pretends to be an 'application', measure its latency
2. Profile & troubleshoot with tools
3. Make changes
4. Repeat

how to measure latency

latency

latency: the time it takes for an operation to complete, such as an application request, a database query, a file system operation, etc.

hiccups



percentiles

Measure percentiles (90th, 99th, 99.9th, ...)

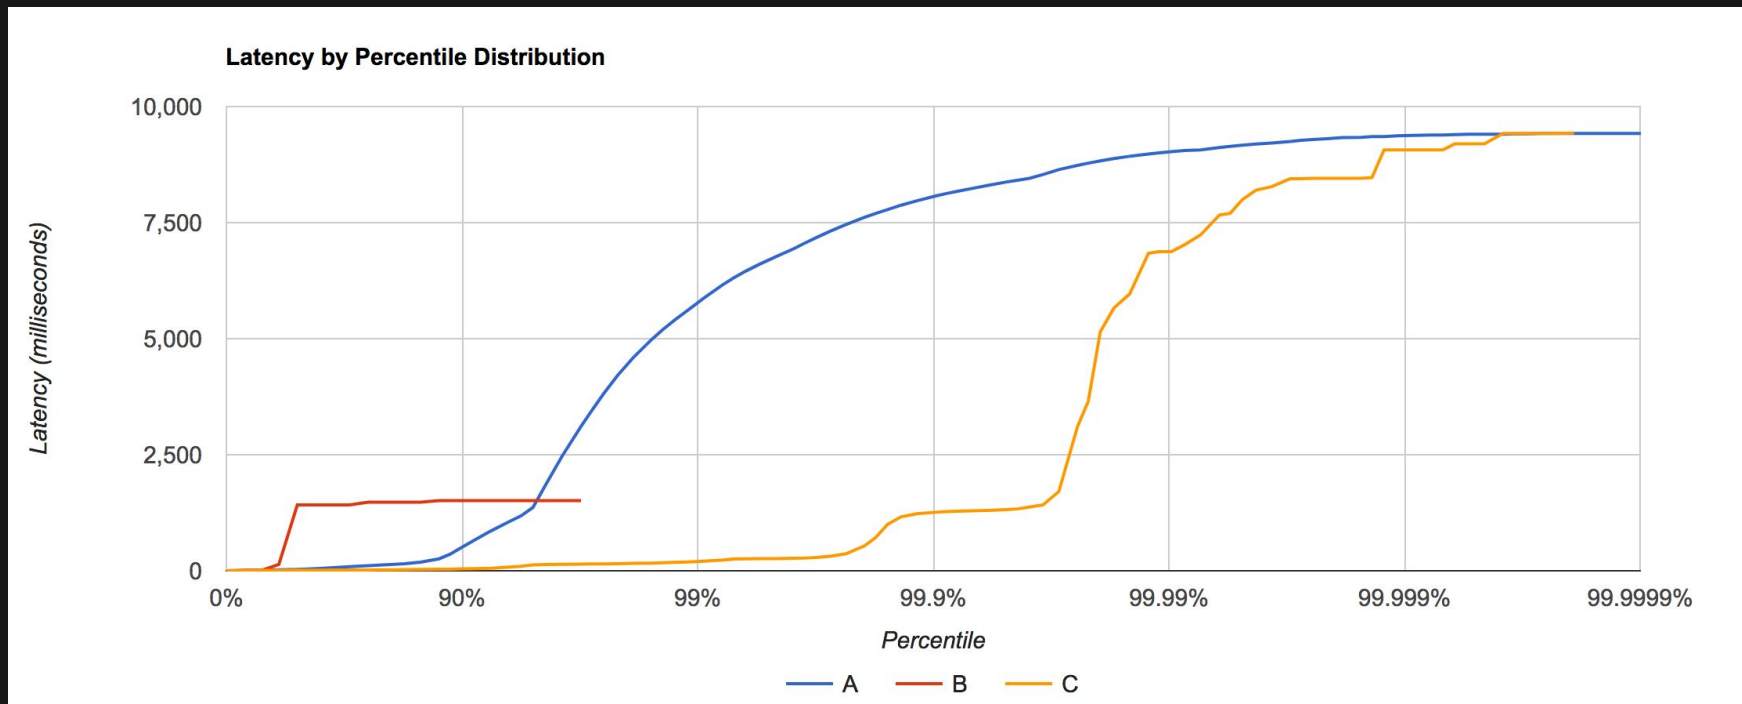
They quantify slowest requests in the population

Think in SLAs: max latency allowed for a percentile

Typically measured as a function of some load

The max is also very meaningful

a latency percentile distribution



coordinated omission

When the loader coordinates with the system under test:

The loader issues requests one by one at a certain rate, but fails to impose that rate due to backoff (e.g., due to synchronous requests);

or

Measure latency before start and end of an operation, but delays outside of timing window do not get measured at all (like queuing and garbage collection pauses).

toolset

toolset

Tools list:

jstack

jmap

gcviewer

gc logs

jit logs

visualvm

gcviewer

verify toolset

```
> ./check_tools.sh  
java ok  
javac ok  
jstat ok  
jvisualvm ok  
jstack ok  
jmap ok  
Java is HotSpot ok  
Java compiler is HotSpot ok  
Java is 1.8 ok  
Java compiler is 1.8 ok
```

how to run an exercise

Take latency measurements (20 secs):

```
> ./run.sh <exercise_number>
```

Latencies in usec recorded to build/Ex<number>-1495666400.hist

Open charts/plotFiles.html in browser, load histogram

For longer profiling / troubleshooting runs (e.g., 5 minutes):

```
> ./run.sh <exercise_number> 300000
```


Exercise 1

A garbage generator

Look at: profiler & GC

Exercise 1

A linked list, storing pairs of strings that get added and removed

1_0

1_0 → base example

Let's execute and see behaviour

```
jstat -gc -h80 $(jps | grep Ex | cut -f1 -d' ') 500
```

S0C	S1C	S0U	S1U	EC	EU	OC	OU	MC	MU	CCSC	CCSU	YGC	YGCT	FGC	FGCT	GCT
0,0	7168,0	0,0	7168,0	57344,0	22528,0	197632,0	47104,0	7552,0	7097,8	896,0	823,3	6	0,199	0	0,000	0,199
0,0	10240,0	0,0	10240,0	74752,0	70656,0	177152,0	117760,0	7552,0	7097,8	896,0	823,3	10	0,432	0	0,000	0,432
0,0	2048,0	0,0	2048,0	11264,0	5120,0	248832,0	227840,0	7552,0	7117,5	896,0	823,3	16	0,754	0	0,000	0,754
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	262144,0	7552,0	7117,5	896,0	823,3	22	0,919	1	0,000	0,919
0,0	6144,0	0,0	6144,0	28672,0	0,0	227328,0	190808,7	7552,0	7110,1	896,0	822,2	24	1,012	1	0,630	1,642
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261976,6	7552,0	7110,1	896,0	822,2	34	1,316	2	0,630	1,946
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261976,6	7552,0	7110,1	896,0	822,2	34	1,316	2	0,630	1,946
0,0	7168,0	0,0	7168,0	39936,0	36864,0	215040,0	173976,2	7552,0	7110,1	896,0	822,2	36	1,372	2	1,548	2,920
0,0	2048,0	0,0	2048,0	9216,0	9216,0	250880,0	250711,2	7552,0	7110,7	896,0	822,2	45	1,688	2	1,548	3,236
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,2	7552,0	7110,7	896,0	822,2	46	1,798	3	1,548	3,346
0,0	0,0	0,0	0,0	57344,0	54272,0	204800,0	157527,4	7552,0	7110,7	896,0	822,2	47	1,798	3	2,404	4,202
0,0	2048,0	0,0	2048,0	10240,0	10240,0	249856,0	249687,4	7552,0	7110,7	896,0	822,2	57	2,115	3	2,404	4,518
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,4	7552,0	7110,7	896,0	822,2	58	2,223	4	2,404	4,627
0,0	3072,0	0,0	3072,0	12288,0	11264,0	246784,0	214359,4	7552,0	7110,7	896,0	822,2	63	2,404	4	3,030	5,434
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,3	7552,0	7110,7	896,0	822,2	70	2,642	5	3,030	5,672
0,0	7168,0	0,0	7168,0	40960,0	37888,0	214016,0	173487,9	7552,0	7110,7	896,0	822,2	72	2,689	5	3,638	6,327
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,4	7552,0	7110,7	896,0	822,2	82	3,062	6	3,638	6,701
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,4	7552,0	7110,7	896,0	822,2	82	3,062	6	3,638	6,701
0,0	2048,0	0,0	2048,0	11264,0	10240,0	248832,0	227159,4	7552,0	7110,7	896,0	822,2	89	3,290	6	4,273	7,563
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,4	7552,0	7110,7	896,0	822,2	94	3,481	7	4,273	7,754
0,0	7168,0	0,0	7168,0	40960,0	37888,0	214016,0	173802,4	7552,0	7110,7	896,0	822,2	96	3,529	7	4,912	8,441
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,4	7552,0	7110,7	896,0	822,2	106	3,879	8	4,912	8,791
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,4	7552,0	7110,7	896,0	822,2	106	3,879	8	4,912	8,791
0,0	2048,0	0,0	2048,0	11264,0	10240,0	248832,0	233815,5	7552,0	7110,7	896,0	822,2	114	4,123	8	5,548	9,671
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,5	7552,0	7110,7	896,0	822,2	118	4,288	9	5,548	9,836
0,0	7168,0	0,0	7168,0	40960,0	37888,0	214016,0	173419,8	7552,0	7110,7	896,0	822,2	120	4,336	9	6,210	10,546
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,3	7552,0	7110,7	896,0	822,2	130	4,693	10	6,210	10,903
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,3	7552,0	7110,7	896,0	822,2	130	4,693	10	6,210	10,903
0,0	2048,0	0,0	2048,0	11264,0	10240,0	248832,0	237911,5	7552,0	7110,7	896,0	822,2	139	4,957	10	6,853	11,809
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,5	7552,0	7110,7	896,0	822,2	143	5,113	11	6,853	11,966
0,0	6144,0	0,0	6144,0	28672,0	8192,0	227328,0	190807,4	7552,0	7110,7	896,0	822,2	145	5,205	11	7,488	12,693
0,0	0,0	0,0	0,0	0,0	0,0	262144,0	261975,4	7552,0	7110,7	896,0	822,2	155	5,517	12	7,488	13,005

```
less ./build/Ex_1_0_gc_$timestamp.log
```

2017-05-25T14:37:27.062-0100: 1,298: [GC pause (G1 Evacuation Pause) (young) (initial-mark), 0,0834642 secs]

[Parallel Time: 83,0 ms, GC Workers: 4]

[GC Worker Start (ms): Min: 1297,9, Avg: 1298,0, Max: 1298,1, Diff: 0,2]

[Ext Root Scanning (ms): Min: 0,1, Avg: 1,2, Max: 4,4, Diff: 4,2, Sum: 5,0]

[Update RS (ms): Min: 0,0, Avg: 2,2, Max: 3,1, Diff: 3,1, Sum: 8,9]

[Processed Buffers: Min: 0, Avg: 3,2, Max: 5, Diff: 5, Sum: 13]

[Scan RS (ms): Min: 0,0, Avg: 0,6, Max: 1,0, Diff: 1,0, Sum: 2,5]

[Code Root Scanning (ms): Min: 0,0, Avg: 0,0, Max: 0,0, Diff: 0,0, Sum: 0,0]

[Object Copy (ms): Min: 78,4, Avg: 78,5, Max: 78,6, Diff: 0,2, Sum: 314,1]

[Termination (ms): Min: 0,1, Avg: 0,2, Max: 0,4, Diff: 0,3, Sum: 0,9]

[Termination Attempts: Min: 21, Avg: 34,0, Max: 49, Diff: 28, Sum: 136]

[GC Worker Other (ms): Min: 0,0, Avg: 0,0, Max: 0,0, Diff: 0,0, Sum: 0,1]

[GC Worker Total (ms): Min: 82,7, Avg: 82,9, Max: 83,0, Diff: 0,2, Sum: 331,4]

[GC Worker End (ms): Min: 1380,8, Avg: 1380,8, Max: 1380,8, Diff: 0,0]

[Code Root Fixup: 0,0 ms]

[Code Root Purge: 0,0 ms]

[Clear CT: 0,1 ms]

[Other: 0,4 ms]

[Choose CSet: 0,0 ms]

[Ref Proc: 0,1 ms]

[Ref Enq: 0,0 ms]

[Redirty Cards: 0,1 ms]

[Humongous Register: 0,0 ms]

[Humongous Reclaim: 0,0 ms]

[Free CSet: 0,0 ms]

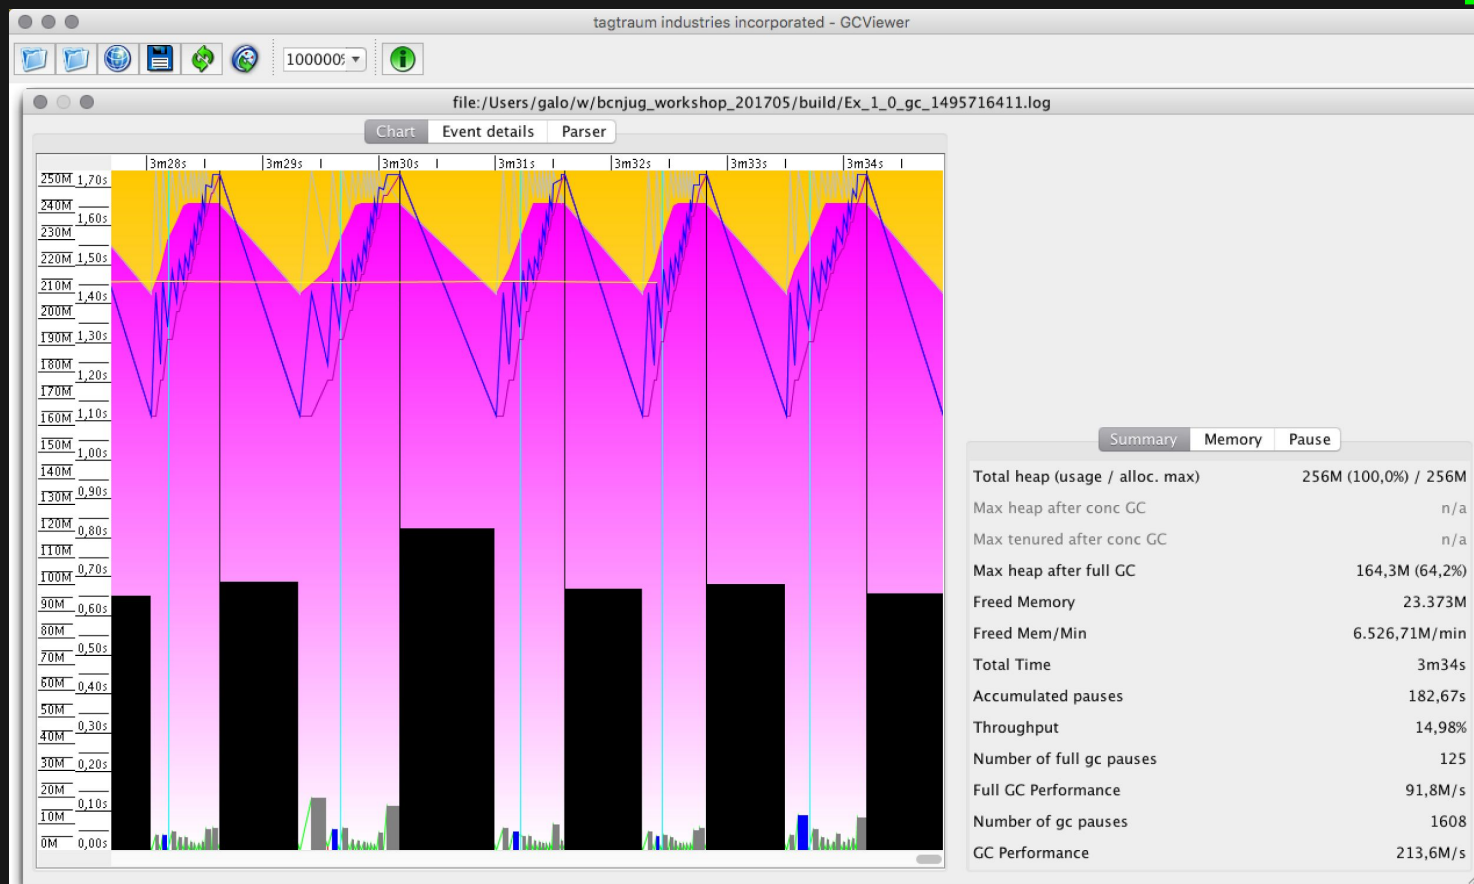
[Eden: 64,0M(64,0M)->0,0B(47,0M) Survivors: 10,0M->10,0M Heap: 199,0M(256,0M)->163,0M(256,0M)]

[Times: user=0,26 sys=0,01, real=0,08 secs]

2017-05-25T14:37:27.146-0100: 1,381: Total time for which application threads were stopped: 0,0836429 seconds, Stopping threads took: 0,0000183 seconds

GCViewer → visualize GC logs

1_0

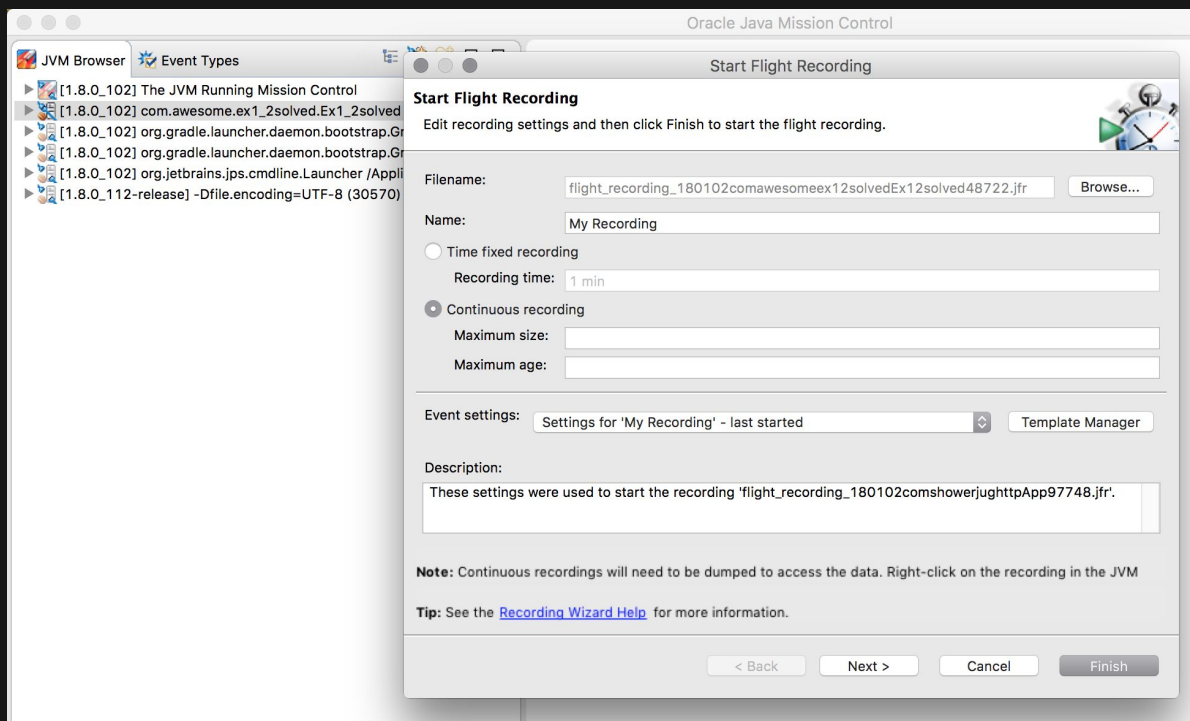


jvisualvm / YourKit → Select process → Monitor + Visual GC

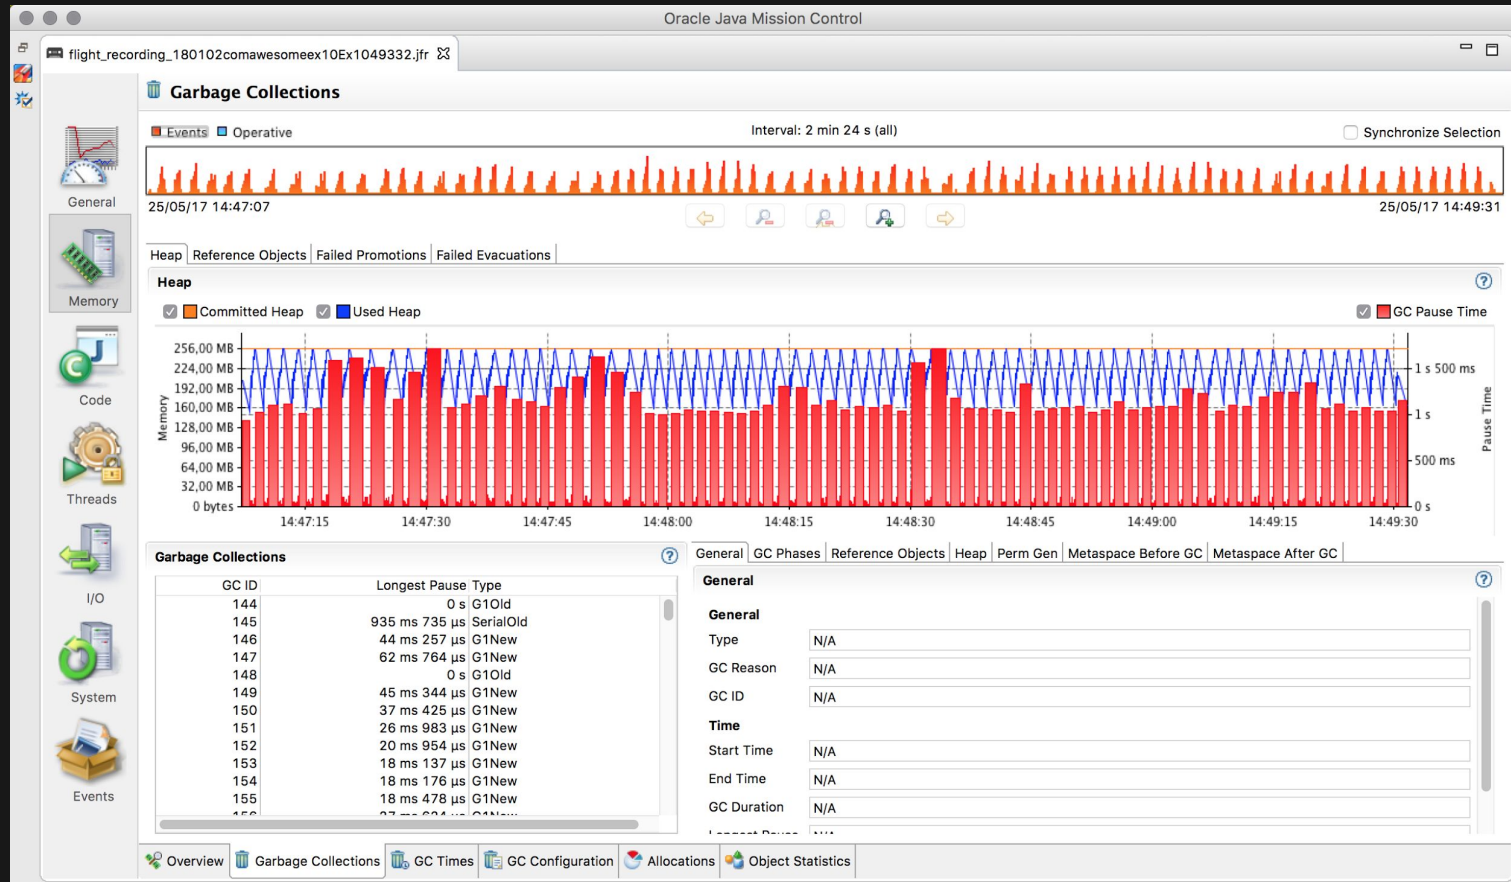
jmc → Select process → Record → Open recording

jmc

→ Select process → Record → open recording



1_0



1_1 → some obvious improvements

Rinse and repeat

1_2 → unimplemented

1_2solved → solutions

1_2solved → more improvements

Rinse and repeat

Exercise 1: Takeaways

JVM is good at dealing with garbage

But you need to help it

Avoid unnecessary copies (e.g. `new String("a")`)

Watch memory overhead of collections & data structures

Exercise 2

A 2-threaded producer / consumer

Look at: profiler & GC

Let's replace the linked list with a specialized concurrent queue:

```
java.util.concurrent.LinkedBlockingQueue
```

2_1 → unimplemented

2_1solved → solution

A non-blocking array-backed queue:

```
org.jctools.queues.SpmcArrayQueue
```

2_2 → unimplemented

2_2solved → solution

Bonus points:

If the queue is full, the producer can become a consumer.

Exercise 2: Takeaways

Coordination is expensive

Don't use synchronized, use specialized data structures

Unbounded queues:

- Produce garbage

- Don't give back pressure, could fill your heap

Non-blocking queues much faster. But you spin on them.

Exercise 3

An operation looks expensive. How can it be improved?

Exercise 3: Takeaways

The Flyweight pattern, while ugly to write, can give a good performance boost as it avoids the object representation

That same technique can be used for off-heap data structures

Slightly better with Scala's value types

Exercise 4

A set of unary operations, which can all execute in a single CPU cycle and generate no garbage, produce an uncommon latency profile.

What's going on?

Hint: you'll probably need to activate some JVM flags in `run.sh`

Exercise 4: (De)Optimizations

The JIT can make a bet, but later revoke it through the use of traps

Optimizations apply at a given call site

A call site can be monomorphic (only one observed type), bimorphic or megamorphic (with different frequencies of observed types)

Class Hierarchy Analysis easily optimizes the mono and bimorphic cases

Exercise 4: Tiered Compilation

```
enum CompLevel {  
    CompLevel_none           = 0, // Interpreter  
    CompLevel_simple         = 1, // C1  
    CompLevel_limited_profile = 2, // C1, invocation & backedge counters  
    CompLevel_full_profile   = 3, // C1, invocation & backedge counters + mdo  
    CompLevel_full_optimization = 4, // C2  
}
```

Exercise 4: Analysis

Running Ex5

Running **loop 1**

com.awesome.ex5.Ex5\$ **DoubleOperation**

8862	90	4	com.awesome.ex5.Ex5::operation (138 bytes)	made
not entrant				

...

8863	101	3	com.awesome.ex5.Ex5::operation (138 bytes)	
------	-----	---	--	--

Exercise 4: Analysis

```
com.awesome.ex5.Ex5$HalfOperation
```

```
  12841   94       4      com.awesome.ex5.Ex5::measureOp (20 bytes)   made  
not entrant
```

```
...
```

```
 12842  103       3      com.awesome.ex5.Ex5::measureOp (20 bytes)
```

```
...
```

```
 12875  110       4      com.awesome.ex5.Ex5::measureOp (20 bytes)
```

```
 12881  103       3      com.awesome.ex5.Ex5::measureOp (20 bytes)   made
```

```
not entrant
```

Exercise 4: Analysis

```
com.awesome.ex5.Ex5$HalfOperation
```

```
  12841    94      4      com.awesome.ex5.Ex5::measureOp (20 bytes)  made  
not entrant
```

```
  <uncommon_trap thread='7171'
```

```
    reason='class_check'
```

```
    action='maybe_recompile'
```

```
    compile_id='94'
```

```
    compiler='C2' level='4'
```

```
    count='3'
```

```
    state='class_check' stamp='12.841'>
```

```
  <jvms bci='7' ... decompiles='1' class_check_traps='3'/>
```

```
</uncommon_trap>
```

Exercise 4: Analysis

```
com.awesome.ex5.Ex5$HalfOperation
```

```
  12841   94       4      com.awesome.ex5.Ex5::measureOp (20 bytes)   made  
not entrant
```

```
...
```

```
  12842  103       3      com.awesome.ex5.Ex5::measureOp (20 bytes)
```

```
...
```

```
  12875  107       4      com.awesome.ex5.Ex5::measureOp (20 bytes)
```

```
  12881  103       3      com.awesome.ex5.Ex5::measureOp (20 bytes)   made
```

```
not entrant
```

Exercise 4: Analysis

```
com.awesome.ex5.Ex5$ IncOperation
```

```
  12224  108      4      com.awesome.ex5.Ex5::measureOp (20 bytes)  made
not entrant
  12224  110      3      com.awesome.ex5.Ex5::measureOp (20 bytes)
  ...
  12227  116      4      com.awesome.ex5.Ex5::measureOp (20 bytes)
  ...
  12232  110      3      com.awesome.ex5.Ex5::measureOp (20 bytes)  made
not entrant
```


Exercise 4: Analysis

```
com.awesome.ex5.Ex5$ DecOperation
```

```
...
```

Exercise 4: Takeaways

Careful with highly polymorphic code in your fast path

JVM performance testing

Many pitfalls:

- Classloading

- Compilation

- Fake warmups

- Garbage Collection

- Eliminated code

Q & A